Working with data: inspect, fix, create

Programming & Data Management, Part II (Python) · Session 2

Magnus Lodefalk
Örebro University

NA446A, Örebro University

Autumn 2026



Today's goals

By the end of today you can:

1. Apply the intake checklist to a variable: type, missing values, range
2. Fix types and missing values, including sentinel codes such as 999
3. Create, subset, sort and aggregate variables, and summarise data in a table and a figure

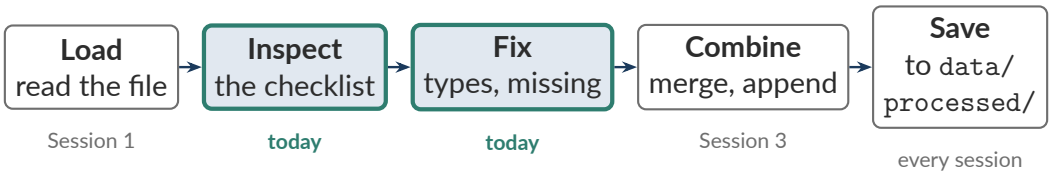
From last time

Answer on your own first, then compare with a neighbour

1. You have cleaned a dataset. Where does your code save it, and why not in `data/raw/`?
2. Write the pandas line that gives mean mpg by origin.

Every data task follows one recipe

Today: inspect and fix, then create and aggregate



- Today's two errors never crash a script: a code read as a number, and a missing value counted as a zero.

Six terms for today

`dtype`	the type of a column: `float64` (number), `int64`, `str` (text)
`NaN`, `<NA>`	pandas' missing value; Stata's `.`
sentinel code	a number that means "missing", such as -9 or 999
boolean mask	a True/False column that picks rows: `cars["mpg"] > 40`
`groupby(...).agg`	one row per group: Stata's `collapse`
`groupby(...).transform`	a group value on every row: Stata's `egen ..., by()`

- A sentinel code looks like data. Only the range check, or the codebook, gives it away.

The intake checklist: ten questions for every variable

Today we practise the four in teal

1. Name, label and unit
2. Type: number, text or date?
3. Missing values, incl. sentinel codes
4. Range: possible min and max?
5. Distribution and outliers
6. Categories: levels and typos
7. Dates: parsed, no gaps
8. IDs: unique, never blank
9. Logic across variables
10. Documentation of every change

– Ask all ten before any regression. Items 7–8 return in Session 3.

▷ The checklist in code

Step 1: load the data and check the types

```
RAW = Path("data/raw")
cars = pd.read_csv(RAW / "auto_messy.csv")
cars.dtypes.head(5)
```

```
mpg             float64
cylinders         int64
displacement    float64
horsepower          str
weight            int64
dtype: object
```

- Horsepower is a number, yet pandas read it as text (`str`): at least one value in the column is not a number.

`auto_messy.csv`: the Session 1 file with two typical problems planted. Stata: `describe`.

Step 2: which values are not numbers?

```
hp = pd.to_numeric(cars["horsepower"],
                   errors="coerce")
cars.loc[hp.isna(), ["name", "horsepower"]]
```

```
                        name horsepower
56     toyota corona hardtop          ?
71           amc matador (sw)         ?
114  chevrolet monte carlo s          ?
151            chevrolet nova         ?
```

`to_numeric` converts text to numbers; `errors="coerce"` turns what it cannot read into `NaN`. `.loc[rows, columns]` shows those rows.

- Four cars have “?” for unknown horsepower. Look before you convert.

Step 3: fix the type

```
cars["horsepower"] = hp
print(cars["horsepower"].dtype)
print(cars["horsepower"].isna().sum())
```

```
float64
4
```

- Horsepower is now a number, and the 4 “?” are real missing values. Count before and after: `coerce` turns anything it cannot read into missing, silently.

Stata: `destring horsepower, replace force` (“4 missing values generated”).

Step 4: count the missing values

```
cars.isna().sum().head(4)
```

```
mpg             0
cylinders       0
displacement    0
horsepower      4
dtype: int64
```

`isna()` is True where a value is missing; `sum()` counts the Trues in each column.

- According to this, mpg has no missing values. Is that true? The next check decides.

Stata: `misstable summarize`, or `count if missing(mpg)`.

Step 5: check the range

```
stats = cars[["mpg", "horsepower"]].describe()
stats.loc[["count", "mean", "min", "max"]].round(1)
```

	mpg	horsepower
count	392.0	388.0
mean	30.9	104.3
min	9.0	46.0
max	999.0	230.0

Ford Pinto (sw), 1972, row 79 mpg 999: one of 3 cars where "missing" was written as a number.

- 999 is a sentinel code, and it pulls the mean up from 23.4 to 30.9.

Step 6: turn the sentinel code into a missing value

```
is_code = cars["mpg"] == 999
print(is_code.sum())
cars["mpg"] = cars["mpg"].mask(is_code)
print(cars["mpg"].isna().sum())
print(cars["mpg"].mean().round(1))
```

```
3
3
23.4
```

`mask(condition)` sets a value to missing where the condition is True.

- 3 codes in, 3 missing out, and the mean falls to 23.4.

Stata: `mvdecode mpg, mv(999)` (“3 missing values generated”).

A car with missing mpg: what does `mpg > 25` give it?

Your bet first: hands up. The answer is on the next slide.

A. 1: a high-mpg car
B. 0: a low-mpg car
C. Missing

Stata says 1, pandas says 0, and both are wrong

Stata: `gen high = mpg > 25`

Missing counts as larger than any number, so the car gets **1**: 157 high-mpg cars

pandas: `(cars["mpg"] > 25).astype(int)`

`NaN > 25` is False, so the car gets **0**: 154 high-mpg cars, and 3 unknowns counted as low

> **Ford Pinto (sw)** Stata calls it a high-mpg car, pandas a low-mpg car. We do not know its mpg.

- Keep missing values missing: Step 8 shows how. Verified in both languages.

Step 7: create variables

```
cars["kpl"] = cars["mpg"] * 0.425144
names = {1: "US", 2: "Europe", 3: "Japan"}
cars["region"] = cars["origin"].map(names)
cars[["mpg", "kpl", "region"]].head(3).round(1)
```

```
    mpg  kpl region
0  18.0  7.7     US
1  15.0  6.4     US
2  18.0  7.7     US
```

- A new column is computed for every row at once, with no loop. `map` turns codes into labels: check them against the codebook (Lab 1's bug).

Stata: `gen kpl = mpg * 0.425144`; `label define` and `label values`. One mile per US gallon is 0.425144 km per litre.

Step 8: a dummy that keeps missing values missing

```
ok = cars["mpg"].notna()
high = (cars["mpg"] > 25).astype("Int8")
cars["high_mpg"] = high.where(ok)
cars["high_mpg"].value_counts(dropna=False)
```

```
high_mpg
0       235
1       154
<NA>      3
Name: count, dtype: Int64
```

- 235 + 154 + 3 = 392: every car counted, and the Pinto stays unknown.

`where(ok)` keeps a value where `ok` is True, missing elsewhere; `Int8` is an integer type that can hold `<NA>`. Stata: `gen high = mpg > 25 if !missing(mpg)`.

Step 9: subset rows and columns, then sort

```
fast = cars.loc[cars["mpg"] > 40,
                ["name", "mpg"]]
fast = fast.sort_values("mpg",
                        ascending=False)
fast.head(3)
```

```
                    name   mpg
320             mazda glc  46.6
327   honda civic 1500 gl  44.6
323  vw rabbit c (diesel)  44.3
```

- A mask picks the rows (8 cars), a list the columns; `cars` is unchanged.

Stata: `gsort -mpg`, `list name mpg if mpg > 40`. pandas returns a sorted copy (assign it); row labels travel with the rows; missing values go last.

Step 10: one row per group with agg

```
by_region = cars.groupby("region").agg(
    n=("mpg", "count"),
    mean_mpg=("mpg", "mean"),
    mean_hp=("horsepower", "mean"))
by_region.round(1)
```

region	n	mean_mpg	mean_hp
Europe	67	27.6	80.6
Japan	79	30.5	79.6
US	243	20.0	118.9

- One row per region, as Stata's `collapse`; the means of Session 1.

▷ count or size?

Step 11: a group value on every row with `transform`

```
grp = cars.groupby("region")["mpg"]
cars["region_mpg"] = grp.transform("mean")
cars["gap"] = cars["mpg"] - cars["region_mpg"]
show = cars[["mpg", "region_mpg", "gap"]]
show.head(3).round(1)
```

```
    mpg  region_mpg  gap
0  18.0        20.0 -2.0
1  15.0        20.0 -5.0
2  18.0        20.0 -2.0
```

- Still 392 rows: each car gets its region's mean, so you can compare the car with its group.

Stata: `bysort origin: egen region_mpg = mean(mpg)`.

Check your understanding

Answer on your own first, then compare with a neighbour

You want each car's horsepower relative to the mean horsepower of cars with the same number of cylinders. `agg` or `transform`? How many rows does the result have?

Step 12: draw the table as a figure

```
m = by_region["mean_mpg"].sort_values()
fig, ax = plt.subplots(figsize=(6, 2.4))
ax.barh(m.index, m, color="0.35")
ax.set_xlabel("Miles per US gallon (mean)")
ax.spines[["top", "right"]].set_visible(False)
```

One figure (`fig`) holds one set of axes (`ax`); every change is a method of `ax`. `barh`: horizontal bars, sorted so the eye reads a ranking. `"0.35"`: a grey that prints well in black and white.

- Draw from the table you already checked, never from numbers typed by hand.

Needs `import matplotlib.pyplot as plt` at the top of the script. Stata: `graph hbar (mean) mpg, over(origin)`.

Step 13: title, source line, save

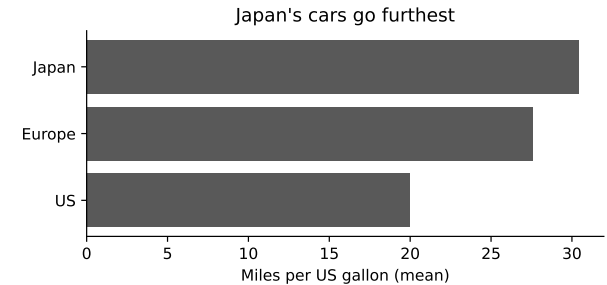
```
ax.set_title("Japan's cars go furthest")
fig.text(0, -0.15, "Source: auto.csv, 389 cars")
for ext in ["png", "pdf"]:
    fig.savefig(f"output/figures/mpg.{ext}",
                bbox_inches="tight")
```

The loop saves the same figure twice: PDF for the paper (sharp at any size), PNG for slides and the web. `bbox_inches="tight"` keeps the source line inside the file.

- Figures go to `output/figures/`, written by code, so a rerun redraws them.

Stata: `graph export "output/figures/mpg.pdf", replace`.

What makes a figure publication-ready



Source: auto.csv, 389 cars

The figure exactly as Step 13 saved it.

1. The title states the finding, not the variable
2. Axes labelled, with units
3. A source line: data and sample
4. Grey that survives black-and-white printing
5. Nothing else: no box, grid, 3D or legend for one series

- If a reader must read the text to understand the figure, the figure is not done.

Today's lab

1. Fill in the TODOs in `s2_lab.py`, one cell at a time *goals 1–3*
2. Check your steps: run `check_s2.py`, fix what it flags *goals 1–3*
3. Ask an AI, then check it: paste your *intake output* (not the data) into a chatbot and ask what problems it shows. Test two of its claims with one line of code each; note claim, test, true or false *rules 3–4*

– A chatbot can read an output well and still be wrong about the data behind it. Your line of code decides.

Never paste personal or register data into a chatbot. Stuck? Use the course tutor prompt (course site, AI page): it gives hints, not answers.

Summary

1. The checklist first: `dtypes`, `isna().sum()`, `describe()` catch a text column and a sentinel code
2. Fix with `to_numeric(errors="coerce")` and `mask`; count missing values before and after every fix
3. Create with `=`, `map` and `where`; pick rows with `.loc`; `agg` gives one row per group, `transform` one value per row; a figure states its finding and its source

Before next time: the Titanic intake exercise (course site), with its self-check; about 90 minutes. Solution after three days.

Test yourself before next time

Without AI: this is what you should manage on your own

1. `income` was read as `str`. Name two likely causes and the line that fixes it.
2. One value was `"1,200"`. What does `pd.to_numeric(..., errors="coerce")` give it, and how would you notice?
3. Write the line that turns the codes -9 and 999 in `age` into missing values.
4. Do `cars["mpg"] <= 25` and `~(cars["mpg"] > 25)` pick the same cars? Why not?
5. Name three things that make a figure publication-ready.

Appendix

The checklist in code

Question	pandas	Stata
1 Name, label, unit	`df.columns`, a codebook file	`label variable`
2 Type	`df.dtypes`, `pd.to_numeric`	`describe`, `destring`
3 Missing, sentinels	`df.isna().sum()`, `mask`	`misstable sum`, `mvdecode`
4 Range	`df.describe()`	`summarize`
5 Distribution, outliers	`df["x"].hist()`, `quantile`	`histogram`, `centile`
6 Categories	`value_counts(dropna=False)`	`tab x, missing`
7 Dates	`pd.to_datetime`	`date()`, `tsset`
8 IDs and keys	`df["id"].is_unique`	`isid id`
9 Cross-variable logic	`assert (df["wage"] > 0).all()`	`assert wage > 0`
10 Documentation	comments, README, codebook	`notes`, do-file

`count` or `size`? And two kinds of missing

- `count` counts non-missing values: Europe has 67 cars with mpg
- `size` counts rows, missing or not: Europe has 68 cars
- Report the one you mean; a table of means should show the `count`

- `NaN` is the missing value of ordinary number columns (`float64`): `NaN > 25` is False
- `<NA>` is the missing value of nullable types (`Int8`, `Float64`): `<NA> > 25` is `<NA>`, so missing stays missing

One more trap: “25 or less” is not “not above 25”

- `len(cars[cars["mpg"] <= 25])` gives 235
- `len(cars[~(cars["mpg"] > 25)])` gives 238: it also keeps the 3 cars with missing mpg, because `NaN > 25` is False and “not False” is True
- In Stata both give 235, since missing counts as greater than 25

- Write the condition you mean, and count the rows it keeps.