Getting data and scaling up

Programming & Data Management, Part II (Python) · Session 4

Magnus Lodefalk
Örebro University

NA446A, Örebro University

Autumn 2026



Today's goals

By the end of today you can:

1. Fetch data from an open API, save the raw answer, and check that every row arrived, as with Stata's `assert _N == ...`
2. Handle dates as with Stata's `monthly()`, `collapse` and `L4.`; write a repeated step once, as Stata's `program`
3. Rebuild everything with one script, as a master do-file, and end with one regression table, as `newey` and `esttab`

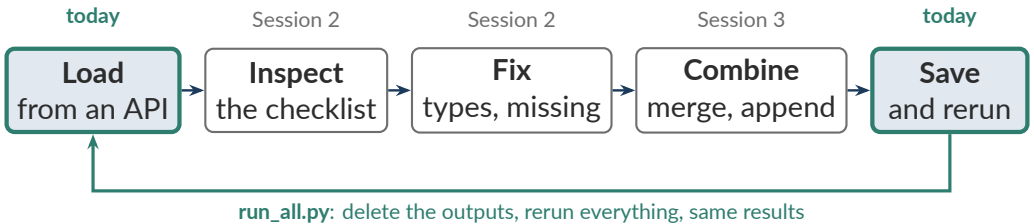
From last time

Answer on your own first, then compare with a neighbour

1. You merge workers onto firms with `validate="m:1"` and pandas stops with a `MergeError`. What does that tell you?
2. Write the line that gives each person's income in the year before, in a panel sorted by person and year.

Every data task follows one recipe

Today: load from an API, then save and rerun (everything)



- Today's two errors never crash a script: an API answer that stops after the first page, and a quarter built from two months.

Today in Stata and in Python

Task	Stata	Python
save a download	`copy "URL" raw/cpi.json`	`requests.post`, `write_text`
text to dates	`gen m = monthly(month, "YM")`	`pd.to_datetime(month, ...)`
months to quarters	`collapse (mean) cpi, by(q)`	`resample("QE").mean()`
yearly growth	`L4.cpi` after `tsset`	`pct_change(4)`
write a step once	`program define`	`def ...: return ...`
rebuild everything	a master do-file	`run_all.py`
regression, table	`newey`, `esttab`	`smf.ols`, HAC, `summary_col`

- Two new words: an API is a web address a program asks for data; JSON is the text it answers in.

Steps 1–6: from SCB's answer to inflation

	Question	Stata you know
1	what does the table hold?	`describe`, before you load
2	how do we fetch the data, and keep a copy?	`copy` the file into `raw/`
3	did every month arrive?	`assert _N == ...`
4	how do month labels become dates?	`monthly()`
5	how do months become quarters?	`collapse (mean) ..., by(q)`
6	how much did prices rise over the year?	`L4.` after `tsset`

- Every step ends with a count: how many rows the source says exist, and how many we have.

Step 1: what does the table hold?

Stata: `describe using cpi` (what a file holds, before you load it)

```
r = requests.get(URL, timeout=30)   # ask SCB what the table holds
r.raise_for_status()                 # stop if SCB said no
meta = r.json()                      # read the answer (JSON)
# 2nd variable: the months
months = meta["variables"][1]["values"]
print(len(months), months[0], months[-1])  # count, first, last
```

```
560 1980M01 2026M08
```

- SCB says the table holds 560 months. Run it today and you get one or two more: SCB adds a month every month.

`.json()`: the text as lists and dictionaries; `["variables"][1]`: the second variable, time. No key needed.

Step 2: how do we fetch the data, and keep a copy?

Stata: `copy "https://..." data/raw/cpi.json` (the raw file first)

```
r = requests.post(URL, json=QUERY, timeout=30)  # ask for the data
r.raise_for_status()                          # stop if SCB said no
(RAW / "cpi.json").write_text(r.text)           # save as it came
rows = r.json()["data"]                         # the data rows
print(rows[0], rows[-1], sep="\n")            # first and last row
```

```
{'key': ['1980M01'], 'values': ['28.38']}
{'key': ['2026M08'], 'values': ['124.85']}
```

- Save the raw answer first: SCB revises its numbers and changes the base year.

Step 3: did every month arrive?

Stata: `assert _N == 560` (561 if you run it today)

```
print(len(rows), len(months))  # rows we got; months SCB lists
# assert _N == 560
assert len(rows) == len(months), "months missing"
```

```
560 560
```

- Compare what you got with what the source says exists. This is the check people skip, because a short answer raises no error.

World Bank, page 1: which countries are on it?

Your bet first: hands up. The answer is on the next slide.

A. All five, about ten rows each
B. All five, the newest years only
C. One country only

Does page 1 hold everything? The World Bank

```
# World Bank: page 1
r = requests.get(WB, params={"format": "json"})
meta_wb, page1 = r.json()          # two parts: the info, the rows
print(meta_wb["total"], meta_wb["pages"])     # rows, pages in all
print(len(page1), page1[0]["countryiso3code"],  # rows on page 1
      page1[-1]["countryiso3code"])       # first and last country
```

```
330 7
50 DNK DNK
```

`WB`: GDP per person for Sweden, Finland, Norway, Denmark and Iceland, 1960–2025.
The answer is a list: first the metadata, then one page of rows.

- 330 rows exist, in 7 pages. Page 1 holds 50, and all 50 are Denmark. A script that reads only page 1 raises no error.

Step 4: how do month labels become dates?

Stata: `gen m = monthly(month, "YM") · destring cpi, replace`

```
cpi = pd.DataFrame(rows)        # rows -> a dataset (Stata: import)
month = cpi["key"].str[-1]            # last part of key: "2026M08"
cpi["date"] = pd.to_datetime(month, format="%YM%m")  # monthly()
cpi["cpi"] = pd.to_numeric(cpi["values"].str[0])     # destring
cpi[["date", "cpi"]].tail(2)                # list date cpi in -2/l
```

```
          date     cpi
558 2026-07-01  125.17
559 2026-08-01  124.85
```

- Now pandas knows these are dates and numbers.

Step 5: how do months become quarters?

Stata: `collapse (mean) cpi (count) months=cpi, by(q)`

```
cpi = cpi.set_index("date")["cpi"]          # tsset date; keep cpi
q = cpi.resample("QE").mean()   # collapse (mean) cpi, by(quarter)
n = cpi.resample("QE").count()   # (count) months=cpi, by(quarter)
pd.DataFrame({"cpi": q, "months": n}).tail(3).round(2)  # list
```

	cpi	months
date		
2026-03-31	124.84	3
2026-06-30	124.84	3
2026-09-30	125.01	2

- Mean for an index or a rate, sum for a flow. The last quarter has 2 months.

`set_index`: dates as row labels; `count()`: months with a value.

Step 6: how much did prices rise over the year?

Stata: `keep if months == 3` · `tsset q` · `gen infl = 100*(cpi/L4.cpi - 1)`

```
q = q[n == 3]                  # keep if months == 3
infl = q.pct_change(4) * 100   # gen infl = 100*(cpi/L4.cpi - 1)
infl.tail(3).round(1)          # list in -3/l
```

```
date
2025-12-31    0.5
2026-03-31    0.5
2026-06-30    0.4
Freq: QE-DEC, Name: cpi, dtype: float64
```

`q[n == 3]`: complete quarters only. `pct_change(4)`: against the same quarter a year before, so seasons cancel.

- Prices rose 0.4 per cent over the year to the second quarter of 2026.

Check your understanding

Answer on your own first, then compare with a neighbour

Predict: (a) you turn monthly *job ads* into quarters with `resample("QE").mean()`. What is wrong? (b) An old answer online writes `resample("Q")`. What happens in pandas 3?

Steps 7–10: from one file to one table

	Question	Stata you know
7	how do we write a repeated step once?	`program define`
8	how do we read every file the same way?	`foreach` (Session 3)
	how do we rebuild everything, and check it?	a master do-file with `log using`
9	does inflation move with unemployment?	`newey infl unemp, lag(4)`
10	how do we save the table for a paper?	`esttab`

- Every step answers one question. Say the question before you write the code.

New Python: a function, as Stata's `program`

Stata

```
program define pct
    args x
    display 100 * `x'
end
pct 0.25
```

Python

```
def pct(x):
    return 100 * x

pct(0.25)
```

```
25.0
```

`def pct(x):`
the colon and the indent

the name and the input, like `program define` and `args`

`return`
the body, like the lines up to Stata's `end`

hands the result back, so you can save it: `y = pct(0.25)`

Step 7: how do we write a repeated step once?

Stata: `program define read_scb` (Steps 2 and 4 inside, then `end`)

```
def read_scb(path):            # program define read_scb; args path
    # read the file's rows
    d = pd.DataFrame(json.loads(path.read_text())["data"])
    month = d["key"].str[-1]                              # "2026M08"
    d.index = pd.to_datetime(month, format="%YM%m")  # monthly()
    return pd.to_numeric(d["values"].str[0])     # destring, return
```

Steps 2 and 4 in one place. Running the cell prints nothing: it only *defines* the function.

- Write a step once, name it, and test it. A fix in the function fixes every file that uses it.

`json.loads`: text to lists and dictionaries. Needs `import json`. The lab's version adds `encoding="utf-8"` to `read_text`, for Windows.

Step 8: how do we read every file the same way?

Stata: `foreach f in cpi unemp { ... }` (the loop from Session 3)

```
series = {}                              # an empty box for the series
for path in sorted(RAW.glob("*.json")):  # foreach f of all .json
    series[path.stem] = read_scb(path)   # read_scb `f'
monthly = pd.concat(series, axis=1)      # merge 1:1 on date
monthly.tail(2)                          # list in -2/l
```

```
               cpi  unemp
key
2026-07-01  125.17    8.6
2026-08-01  124.85    8.9
```

- Two files today, fifty tomorrow: the loop does not change.

How do we rebuild everything? `run_all.py`

Stata: `master.do: log using run_log, then one do per script, in order`

```
SCRIPTS = ["scripts/01_get_data.py", "scripts/02_clean.py",
           "scripts/03_table.py"]
```

```
$ uv run run_all.py
01_get_data.py: cpi 560/560 months, unemp 308/308 months (SCB API)
02_clean.py: 868 monthly values in, 102 quarters out (2001Q1-2026Q2)
  dropped 84 of 186 quarters with a gap (84 before the unemployment data)
03_table.py: 102 quarters in, 2 models out (output/tables/regression.*)
run_all done: log in output/logs/run_log.txt
```

- The list is the recipe; each script prints rows in, out, and why.

The delete-outputs test

Gentzkow and Shapiro (2014), *Code and Data for the Social Sciences*, ch. 2

1. Delete `data/processed/` and `output/`: everything the code made
2. Run `run_all.py` once
3. Everything comes back, with the same numbers; the log in `output/logs/run_log.txt` shows rows in and out of every step

   "The proof in the pudding is that we can now delete all of the output files in the directory [...] and reproduce them by running `rundirectory.bat`. This is the precise sense in which the output is now replicable."

– The log is the first thing to read when a number looks strange: a step that drops half the sample shows there before it shows in the regression.

▷ The same pipeline in Stata

Step 9: does inflation move with unemployment?

Stata: `newey infl unemp, lag(4)` · `newey infl unemp L.infl, lag(4)`

```
q = pd.read_parquet("data/processed/quarterly.parquet")  # use
m1 = smf.ols("infl ~ unemp", data=q).fit(     # regress infl unemp
    cov_type="HAC", cov_kwds={"maxlags": 4})       # newey, lag(4)
# adds L.infl
m2 = smf.ols("infl ~ unemp + infl_lag1", data=q).fit(
    cov_type="HAC", cov_kwds={"maxlags": 4})       # newey, lag(4)
```

A Phillips-curve regression: quarterly inflation on the unemployment rate, 2001–2026; model 2 adds last quarter's inflation. HAC (Newey–West) standard errors allow errors correlated over time.

- The regression reads the file the pipeline saved, never numbers typed by hand.

`"infl ~ unemp"`: inflation on unemployment, constant added, as `regress infl unemp`. Stata's `newey` gives the same coefficients.

Step 10: how do we save the table for a paper?

Stata: `eststo m1: newey ... · esttab m1 m2 using regression.csv, se`

```
t = summary_col([m1, m2], stars=True,          # esttab m1 m2, star
                float_format="%.2f")           # b(%9.2f) se(%9.2f)
# add the N row
t.tables[0].loc["N"] = [str(int(m1.nobs)), str(int(m2.nobs))]
t.tables[0].to_csv(TABLES / "regression.csv")  # esttab using .csv
(TABLES / "regression.tex").write_text(t.as_latex())  # and .tex
```

```
607
```

`summary_col`: models side by side, as `esttab`. Line 3 adds each model's N. The 607 is the number of characters `write_text` wrote.

- The CSV opens in Excel and Word; the `.tex` goes straight into a paper. Never copy numbers from the screen.

Needs `from statsmodels.iolib.summary2 import summary_col`. Stata's `esttab` and `eststo` come from SSC.

The table, as the pipeline saved it

```
================================
                 infl I  infl II
--------------------------------
Intercept       6.46***  0.98
                (1.35)   (0.79)
unemp           -0.61*** -0.11
                (0.18)   (0.10)
infl_lag1                0.93***
                         (0.09)
R-squared       0.06     0.88
R-squared Adj.  0.05     0.88
N               102      102
================================
```

- Model I: a quarter with one point higher unemployment had 0.6 points lower inflation, on average
- Model II: most of that is last quarter's inflation carrying over (0.93); the unemployment coefficient falls to -0.11
- A correlation over time, not a causal effect

- N = 102 is the number in the log. The table and the log must agree.

▷ Rerun next month

For the exam period: predict, test, explain

1. Predict before you run: write down what the line prints. The exam asks exactly that, in WISEflow, without Python and without AI
2. Test on a case whose answer you know: count rows in and out, `assert` what must hold, compute one number by hand
3. Explain every line in one sentence, including the lines an AI wrote. If you cannot, you do not own the result yet

– Code that runs is not code that is right. Code you can predict, test and explain is.

Practise: the sample exam and the practice bank on the course site, without AI.

Today's lab: a mini-pipeline

1. Unzip `s4-lab.zip` into `na446a-starter`; open `s4-lab/scripts/01_get_data.py`; run a script whole with ▷
2. In class: TODO 1, then 2a–2d in `02_clean.py`. Each TODO shows the Stata line; you write the Python
3. If time: review an AI's pipeline, `ai_02_clean.py`: three `asserts` that catch it, then fix it
4. At home: TODOs 3–5, then `check_s4.py` (it does the delete-outputs test)

– No network, or SCB answers 429 ("too many requests")? Set `USE_SAVED = True`: the saved copy runs the same way.

Never paste personal or register data into a chatbot. API keys (none today) go in `.env`, never in code.

Summary

1. An API: ask for the metadata, ask for the data, save the raw answer, `assert` the count; an answer in pages needs a loop to the last page
2. Dates: `to_datetime` with a format; `resample("QE")` with a rule and a count; `pct_change(4)`; one function, one loop over the files
3. `run_all.py` rebuilds everything and logs rows in and out, as a master do-file; `summary_col` gives one table as CSV and LaTeX, as `esttab`

Self-study: DuckDB, which reads only what a query needs, so a file need not fit in memory (appendix).
Before the exam: the World Bank exercise and its self-check (about 90 minutes; solution after three days), then the sample exam, without AI.

Test yourself before next time

Without AI: this is what you should manage on your own

1. The metadata lists 308 months; your data have 300 rows. Name two likely causes and the line that stops your script.
2. Write the line that turns labels such as `"2024M07"` into dates.
3. Monthly hours worked to quarters: which rule, and why? And for the unemployment rate?
4. A function lacks `return`: what does it give back, and `pd.concat` then?
5. Write the SQL that gives the mean of `infl` per year from `quarterly.parquet`.
6. What is the delete-outputs test, and what does it prove?
7. A script reads only the World Bank's page 1. What does it report, and what stops it?

Appendix

The question SCB receives (QUERY)

```
QUERY = {"query": [                    # what we ask SCB for:
    {"code": "ContentsCode",           # the variable: CPI ...
     "selection": {"filter": "item", "values": ["00000807"]}},
    {"code": "Tid",                    # ... for every month ("*")
     "selection": {"filter": "all", "values": ["*"]}}],
    "response": {"format": "json"}}  # answer in JSON
```

- One entry per variable of the table: here the series (`00000807`, the CPI since 1980) and time (`Tid`, all months: `"*"`)
- A table with more variables needs one entry each: the lab's unemployment rate picks labour status, seasonal adjustment, sex and age
- The lab's `01_get_data.py` builds the query in a loop

Every page, with a loop (the exercise)

```
meta, rows = get_page(name, 1)  # page 1, and the info
# forvalues p = 2/`pages'
for page in range(2, meta["pages"] + 1):
    rows = rows + get_page(name, page)[1]  # append it
# assert _N == total
assert len(rows) == meta["total"], "rows missing"
```

- Page 1 says how many pages exist; the loop asks for the rest
- `get_page` (given in the exercise) sends `page=...` with the request and saves each page in `data/raw/wb/`
- Last year's course code set `per_page=130` for 132 rows and silently lost two years of Swedish data: a larger page is not a check

Self-study: DuckDB, for a file too large to load

```
duckdb.sql("""
  SELECT year(date) AS year, AVG(infl) AS infl
  FROM 'data/processed/quarterly.parquet'
  WHERE year(date) >= 2024
  GROUP BY year ORDER BY year""").df().round(2)
```

	year	infl
0	2024	2.86
1	2025	0.68
2	2026	0.47

- DuckDB reads only the rows and columns it needs, so the file need not fit in memory.

SQL: the language of databases. Stata has nothing like it: `use` loads the whole file.

The same pipeline in Stata (run on the lab's data)

pandas	Stata
`pd.to_datetime(month, format="%YM%m")`	`gen m = monthly(month, "YM")`
`resample("QE").mean(), .count()`	`collapse (mean) cpi (count) n=cpi, by(q)`
`pct_change(4) * 100`	`tsset q` then `gen infl = 100*(cpi/L4.cpi - 1)`
`shift(1)`	`gen infl_lag1 = L.infl`
`def read_scb(path): ...`	`program define`
`for path in sorted(RAW.glob("*.json")):`	`foreach f of local files { ... }`
`run_all.py` with a log	a master do-file with `log using`
`fit(cov_type="HAC", ...)`	`newey infl unemp, lag(4)`
`summary_col, as_latex()`	`esttab` (from SSC)

Same 102 quarters and the same coefficients: -0.61 (0.18); with the lag -0.11 and 0.93. The intercept's standard error differs in the second decimal (1.35 vs 1.37): statsmodels applies a small-sample correction.

API keys go in `.env`, never in code

```
from dotenv import load_dotenv
import os                          # talks to the computer
load_dotenv()                    # reads .env
key = os.getenv("FRED_API_KEY")  # the key, never in the script
```

- `.env` is a text file in the project folder with one line: `FRED_API_KEY=...`
- `.gitignore` lists `.env`, so the key never reaches GitHub
- SCB and the World Bank need no key; FRED (US data) does

Rerun next month: what changes?

- After SCB publishes September for both series, the third quarter of 2026 has three months: N rises from 102 to 103, and every coefficient moves a little
- The raw answers change too: SCB revises its numbers and changes the base year
- To reproduce today's table later, keep today's saved answers: `USE_SAVED = True` uses the copy of 29 September
- A pipeline that reruns is not a result that stays: the saved raw answer is what makes a table reproducible.